

Fraud Detection in Financial Transactions

1. Introduction

This project focuses on detecting fraudulent financial transactions using machine learning techniques. Fraud detection is crucial in banking systems to prevent financial loss and ensure security. In today's digital era, financial transactions are increasingly conducted online through banking systems, credit cards, mobile wallets, and e-commerce platforms. While this digital transformation has improved convenience and speed, it has also led to a significant rise in fraudulent activities. Fraudulent transactions can cause substantial financial losses to individuals, businesses, and financial institutions, making fraud detection a critical area of concern.

Fraud detection involves identifying unusual or suspicious patterns in transaction data that may indicate illegal or unauthorized activities. Traditional rule-based systems are often insufficient because fraudsters continuously evolve their techniques to bypass security measures. Therefore, advanced techniques such as Machine Learning (ML) have become essential for detecting fraud effectively.

Machine learning models can analyze large volumes of transaction data, learn patterns of normal behavior, and identify anomalies that deviate from these patterns. In this project, we use an **anomaly detection approach**, specifically the Isolation Forest algorithm, to detect fraudulent transactions. This method is particularly useful when fraud cases are rare compared to normal transactions, which is a common scenario in real-world datasets.

The dataset used in this project contains transaction-related features such as transaction amount, location, merchant details, age, and gender. The goal is to build a model that can accurately distinguish between legitimate and fraudulent transactions.

This project not only demonstrates the application of machine learning in fraud detection but also highlights the importance of data preprocessing, handling class imbalance, and evaluating model performance using metrics such as confusion matrix, ROC curve, and precision-recall curve.

2. Dataset Description

The dataset contains the following columns:

- transaction_id
- transaction_amount
- Location
- Merchant
- Age
- Gender
- Fraud (0 = No, 1 = Yes)

3. Objective

To identify suspicious or fraudulent transactions using anomaly detection techniques. The main objectives of this project are:

- 1. To detect fraudulent transactions**
Identify suspicious or unauthorized financial transactions using machine learning techniques.
- 2. To analyze transaction data**
Study patterns in transaction attributes such as amount, location, merchant, age, and gender to understand normal vs abnormal behavior.
- 3. To apply anomaly detection techniques**
Use algorithms like Isolation Forest to detect rare and unusual transactions that may indicate fraud.
- 4. To handle class imbalance**
Address the issue where fraudulent transactions are much fewer than normal transactions, ensuring the model performs effectively.
- 5. To evaluate model performance**
Measure the effectiveness of the model using metrics such as confusion matrix, precision, recall, F1-score, ROC curve, and precision-recall curve.
- 6. To visualize results clearly**
Create graphs and plots to compare actual and predicted fraud cases for better understanding and interpretation.
- 7. To develop a simple fraud alert system**
Highlight transactions that are predicted as fraudulent for further investigation.

4. Methodology

The methodology of this project explains the step-by-step process followed to detect fraudulent transactions using machine learning techniques.

Step 1: Data Collection

The dataset containing financial transaction details is collected. It includes features such as transaction ID, transaction amount, location, merchant name, age, gender, and fraud label (0 = normal, 1 = fraud).

Step 2: Data Preprocessing

Data preprocessing is performed to clean and prepare the dataset for modeling:

- Missing values are removed or handled.
- Irrelevant columns such as transaction ID are dropped.
- Categorical variables (Location, Merchant, Gender) are converted into numerical format using encoding techniques.
- Data consistency and correctness are ensured.

Step 3: Feature Scaling

Since machine learning models perform better with scaled data, numerical features are standardized using techniques like StandardScaler. This ensures all features contribute equally to the model.

Step 4: Exploratory Data Analysis (EDA)

Basic analysis is performed to understand the dataset:

- Distribution of fraud vs non-fraud transactions
- Relationship between transaction amount and fraud
- Correlation between features using heatmaps

Step 5: Handling Class Imbalance

In fraud detection datasets, fraudulent transactions are usually very few compared to normal ones. This imbalance is considered while building the model to avoid biased predictions.

Step 6: Model Selection

An anomaly detection algorithm, **Isolation Forest**, is selected because:

- It works well for detecting rare events (fraud)
- It does not require balanced data
- It isolates anomalies effectively

Step 7: Model Training

The Isolation Forest model is trained on the processed dataset. It learns patterns of normal transactions and identifies deviations as anomalies.

Step 8: Prediction

The trained model predicts whether each transaction is normal or fraudulent:

- Output 1 → Normal
- Output -1 → Anomaly (converted to Fraud = 1)

Step 9: Model Evaluation

The performance of the model is evaluated using:

- Confusion Matrix
- Precision, Recall, F1-score
- ROC Curve
- Precision-Recall Curve

Step 10: Visualization

Graphs and plots are created for better understanding:

- Confusion matrix heatmap
- Actual vs predicted comparison

- ROC curve
- Transaction amount vs fraud

Step 11: Result Generation

The predicted results are stored in a new dataset with a column **Predicted_Fraud**. Fraudulent transactions are highlighted for further investigation.


Step 12: Fraud Alert System

A simple alert system is created to display transactions that are predicted as fraudulent, helping in quick decision-making.

5. Visualizations

• Feature Correlation Heatmap

A feature correlation heatmap visually represents the relationship between different variables in the dataset using correlation values ranging from -1 to +1. Positive values indicate a direct relationship, while negative values indicate an inverse relationship. In this heatmap, transaction_amount shows a strong positive correlation with fraud, making it an important feature for prediction, whereas most other features like age and transaction_id have weak correlations, indicating less influence. It helps in understanding which features are most relevant and how they are related to each other.




Interpretation :

The correlation heatmap shows that transaction_amount has a strong positive correlation with fraud (0.77), indicating higher transactions are more likely to be fraudulent. Merchant and gender are highly correlated (0.91), suggesting a strong relationship between these features. Location has a strong negative correlation with merchant and gender (-0.77), showing an inverse relationship. Other features like age and transaction_id have very weak correlations with fraud, meaning they have less impact on fraud prediction. Overall, transaction amount is the most important factor influencing fraud in this dataset.

Confusion matrix Heatmap:

A confusion matrix is a table used to evaluate a model by comparing actual and predicted values, showing correct and incorrect predictions (TP, TN, FP, FN) to measure performance.



Interpretation :

The confusion matrix shows that the model correctly identified 67 normal transactions (True Negatives) and 2 fraud cases (True Positives). However, it incorrectly classified 3 normal transactions as fraud (False Positives) and missed 14 fraud cases (False Negatives). This indicates that while the model performs well in detecting normal transactions, it struggles to identify fraud cases, missing a significant number of them, which is a critical issue in fraud detection.


Classification Report:

Classification Report:				
	precision	recall	f1-score	support
0	0.83	0.96	0.89	70
1	0.40	0.12	0.19	16
accuracy			0.80	86
macro avg	0.61	0.54	0.54	86
weighted avg	0.75	0.80	0.76	86

Interpretation:

The model accuracy is about 80%, which means it correctly predicts 80 out of 100 transactions overall. However, this can be misleading in fraud detection because the model is much better at identifying normal transactions than fraud cases. As seen in the confusion matrix, it misses many fraud cases (high false negatives), so even with good accuracy, the model’s performance on detecting fraud is weak.

• Actual vs Predicted Comparison



Interpretation:

The chart shows that the model is biased toward predicting the majority class (normal transactions), overestimating normal cases while significantly underestimating fraud cases. Although it may achieve decent overall accuracy, it performs poorly at detecting fraud, missing many actual fraud instances (high false negatives). This indicates a class imbalance issue and suggests the model needs improvement, especially in recall for the fraud class, to be reliable in practice.

6. Results

The fraud detection model achieved an accuracy of around 80%, correctly identifying most of the normal transactions. The confusion matrix shows that the model performs well in predicting non-fraud cases but struggles with detecting fraudulent transactions, as several fraud cases are misclassified as normal. The analysis also revealed that transaction amount is the most influential feature in identifying fraud. Visualization techniques such as the confusion matrix, and correlation heatmap helped in understanding model performance and feature relationships.

7. Conclusion

Fraud detection using machine learning helps in identifying suspicious activities efficiently. Further improvements can be made using deep learning models and larger datasets. This project demonstrates that machine learning techniques, particularly anomaly detection using Isolation Forest, can be effectively applied to detect fraudulent financial transactions. While the model shows good overall accuracy, its ability to detect fraud is limited due to class imbalance and the small dataset size. Therefore, improving fraud detection requires more data, better handling of imbalance, and possibly advanced models like deep learning or ensemble methods. Overall, the project highlights the importance of accurate fraud detection systems in financial security and provides a strong foundation for further improvement.